

# Overview of x64 ABI conventions

Article • 04/22/2022

This topic describes the basic application binary interface (ABI) for x64, the 64-bit extension to the x86 architecture. It covers topics such as the calling convention, type layout, stack and register usage, and more.

## x64 calling conventions

Two important differences between x86 and x64 are:

- 64-bit addressing capability
- Sixteen 64-bit registers for general use.

Given the expanded register set, x64 uses the `__fastcall` calling convention and a RISC-based exception-handling model.

The `__fastcall` convention uses registers for the first four arguments, and the stack frame to pass more arguments. For details on the x64 calling convention, including register usage, stack parameters, return values, and stack unwinding, see [x64 calling convention](#).

For more information on the `__vectorcall` calling convention, see [\\_\\_vectorcall](#).

## Enable x64 compiler optimization

The following compiler option helps you optimize your application for x64:

- [/favor \(Optimize for Architecture Specifics\)](#)


## x64 type and storage layout

This section describes the storage of data types for the x64 architecture.

### Scalar types

Although it's possible to access data with any alignment, align data on its natural boundary, or a multiple of its natural boundary, to avoid performance loss. Enums are constant integers and are treated as 32-bit integers. The following table describes the type definition and recommended storage for data as it pertains to alignment using the following alignment values:

- Byte - 8 bits
- Word - 16 bits
- Doubleword - 32 bits
- Quadword - 64 bits
- Octaword - 128 bits

 Expand table

Scalar type	C data type	Storage size (in bytes)	Recommended alignment
INT8	char	1	Byte
UINT8	unsigned char	1	Byte
INT16	short	2	Word
UINT16	unsigned short	2	Word
INT32	int, long	4	Doubleword
UINT32	unsigned int, unsigned long	4	Doubleword
INT64	__int64	8	Quadword
UINT64	unsigned __int64	8	Quadword
FP32 (single precision)	float	4	Doubleword
FP64 (double precision)	double	8	Quadword
POINTER	*	8	Quadword
__m64	struct __m64	8	Quadword
__m128	struct __m128	16	Octaword

## x64 aggregate and union layout

Other types, such as arrays, structs, and unions, have stricter alignment requirements that ensure consistent aggregate and union storage and data retrieval. Here are the definitions for array, structure, and union:

- Array

Contains an ordered group of adjacent data objects. Each object is called an *element*. All elements within an array have the same size and data type.
- Structure

Contains an ordered group of data objects. Unlike the elements of an array, the members of a structure can have different data types and sizes.
- Union

An object that holds any one of a set of named members. The members of the named set can be of any type. The storage allocated for a union is equal to the storage required for the largest member of that union, plus any padding required for alignment.

The following table shows the strongly recommended alignment for the scalar members of unions and structures.

 Expand table

Scalar Type	C Data Type	Required Alignment
INT8	char	Byte
UINT8	unsigned char	Byte

Scalar Type	C Data Type	Required Alignment
INT16	short	Word
UINT16	unsigned short	Word
INT32	int, long	Doubleword
UINT32	unsigned int, unsigned long	Doubleword
INT64	__int64	Quadword
UINT64	unsigned __int64	Quadword
FP32 (single precision)	float	Doubleword
FP64 (double precision)	double	Quadword
POINTER	*	Quadword
__m64	struct __m64	Quadword
__m128	struct __m128	Octaword

The following aggregate alignment rules apply:

- The alignment of an array is the same as the alignment of one of the elements of the array.
- The alignment of the beginning of a structure or a union is the maximum alignment of any individual member. Each member within the structure or union must be placed at its proper alignment as defined in the previous table, which may require implicit internal padding, depending on the previous member.
- Structure size must be an integral multiple of its alignment, which may require padding after the last member. Since structures and unions can be grouped in arrays, each array element of a structure or union must begin and end at the proper alignment previously determined.
- It's possible to align data in such a way as to be greater than the alignment requirements as long as the previous rules are maintained.
- An individual compiler may adjust the packing of a structure for size reasons. For example, [/Zp \(Struct Member Alignment\)](#) allows for adjusting the packing of structures.

## x64 structure alignment examples


The following four examples each declare an aligned structure or union, and the corresponding figures illustrate the layout of that structure or union in memory. Each column in a figure represents a byte of memory, and the number in the column indicates the displacement of that byte. The name in the second row of each figure corresponds to the name of a variable in the declaration. The shaded columns indicate padding that is required to achieve the specified alignment.

### Example 1

C

```
// Total size = 2 bytes, alignment = 2 bytes (word).
_declspec(align(2)) struct {
```

```
    short a;    // +0; size = 2 bytes
}
```




Example 2

```
C

// Total size = 24 bytes, alignment = 8 bytes (quadword).

_declspec(aligned(8)) struct {
    int a;    // +0; size = 4 bytes
    double b; // +8; size = 8 bytes
    short c;  // +16; size = 2 bytes
}
```



Example 3

```
C

// Total size = 12 bytes, alignment = 4 bytes (doubleword).

_declspec(aligned(4)) struct {
    char a;    // +0; size = 1 byte
    short b;   // +2; size = 2 bytes
    char c;    // +4; size = 1 byte
    int d;     // +8; size = 4 bytes
}
```




Example 4

```
C

// Total size = 8 bytes, alignment = 8 bytes (quadword).

_declspec(aligned(8)) union {
    char *p;    // +0; size = 8 bytes
    short s;    // +0; size = 2 bytes
    long l;     // +0; size = 4 bytes
}
```



## Bitfields

Structure bit fields are limited to 64 bits and can be of type signed int, unsigned int, int64, or unsigned int64. Bit fields that cross the type boundary will skip bits to align the bitfield to the next type alignment. For example, integer bitfields may not cross a 32-bit boundary.

## Conflicts with the x86 compiler

Data types that are larger than 4 bytes aren't automatically aligned on the stack when you use the x86 compiler to compile an application. Because the architecture for the x86 compiler is a 4 byte aligned stack, anything larger than 4 bytes, for example, a 64-bit integer, can't be automatically aligned to an 8-byte address.

Working with unaligned data has two implications.

- It may take longer to access unaligned locations than it takes to access aligned locations.
- Unaligned locations can't be used in interlocked operations.

If you require more strict alignment, use `__declspec(align(N))` on your variable declarations. This causes the compiler to dynamically align the stack to meet your specifications. However, dynamically adjusting the stack at run time may cause slower execution of your application.

## x64 register usage

The x64 architecture provides for 16 general-purpose registers (hereafter referred to as integer registers) as well as 16 XMM/YMM registers available for floating-point use. Volatile registers are scratch registers presumed by the caller to be destroyed across a call. Nonvolatile registers are required to retain their values across a function call and must be saved by the callee if used.

## Register volatility and preservation

The following table describes how each register is used across function calls:

 Expand table

Register	Status	Use
RAX	Volatile	Return value register
RCX	Volatile	First integer argument
RDX	Volatile	Second integer argument
R8	Volatile	Third integer argument
R9	Volatile	Fourth integer argument

Register	Status	Use
R10:R11	Volatile	Must be preserved as needed by caller; used in syscall/sysret instructions
R12:R15	Nonvolatile	Must be preserved by callee
RDI	Nonvolatile	Must be preserved by callee
RSI	Nonvolatile	Must be preserved by callee
RBX	Nonvolatile	Must be preserved by callee
RBP	Nonvolatile	May be used as a frame pointer; must be preserved by callee
RSP	Nonvolatile	Stack pointer
XMM0, YMM0	Volatile	First FP argument; first vector-type argument when <code>__vectorcall</code> is used
XMM1, YMM1	Volatile	Second FP argument; second vector-type argument when <code>__vectorcall</code> is used
XMM2, YMM2	Volatile	Third FP argument; third vector-type argument when <code>__vectorcall</code> is used
XMM3, YMM3	Volatile	Fourth FP argument; fourth vector-type argument when <code>__vectorcall</code> is used
XMM4, YMM4	Volatile	Must be preserved as needed by caller; fifth vector-type argument when <code>__vectorcall</code> is used
XMM5, YMM5	Volatile	Must be preserved as needed by caller; sixth vector-type argument when <code>__vectorcall</code> is used
XMM6:XMM15, YMM6:YMM15	Nonvolatile (XMM), Volatile (upper half of YMM)	Must be preserved by callee. YMM registers must be preserved as needed by caller.

On function exit and on function entry to C Runtime Library calls and Windows system calls, the direction flag in the CPU flags register is expected to be cleared.

## Stack usage

For details on stack allocation, alignment, function types and stack frames on x64, see [x64 stack usage](#).

## Prolog and epilog

Every function that allocates stack space, calls other functions, saves nonvolatile registers, or uses exception handling must have a prolog whose address limits are described in the unwind data associated with the respective function table entry, and epilogs at each exit to a function. For details on the required prolog and epilog code on x64, see [x64 prolog and epilog](#).

## x64 exception handling

For information on the conventions and data structures used to implement structured exception handling and C++ exception handling behavior on the x64, see [x64 exception handling](#).

## Intrinsics and inline assembly

One of the constraints for the x64 compiler is no inline assembler support. This means that functions that can't be written in C or C++ will either have to be written as subroutines or as intrinsic functions supported by the compiler. Certain functions are performance sensitive while others aren't. Performance-sensitive functions should be implemented as intrinsic functions.

The intrinsics supported by the compiler are described in [Compiler intrinsics](#).

## x64 image format

The x64 executable image format is PE32+. Executable images (both DLLs and EXEs) are restricted to a maximum size of 2 gigabytes, so relative addressing with a 32-bit displacement can be used to address static image data. This data includes the import address table, string constants, static global data, and so on.

## See also

[Calling Conventions](#)

---

## Feedback

Was this page helpful?



[Provide product feedback](#)  | [Get help at Microsoft Q&A](#)